



DRC-Bench

Success rate by model and agent framework



Mean over five runs per task. Higher is better; panels use separate scales.

Source: arxiv.org/abs/2605.06936v4 · Main comparison; ablations excluded.

Color identifies the model; the second label line identifies the framework.

ORFS-Agent is PPA-only and is omitted here.